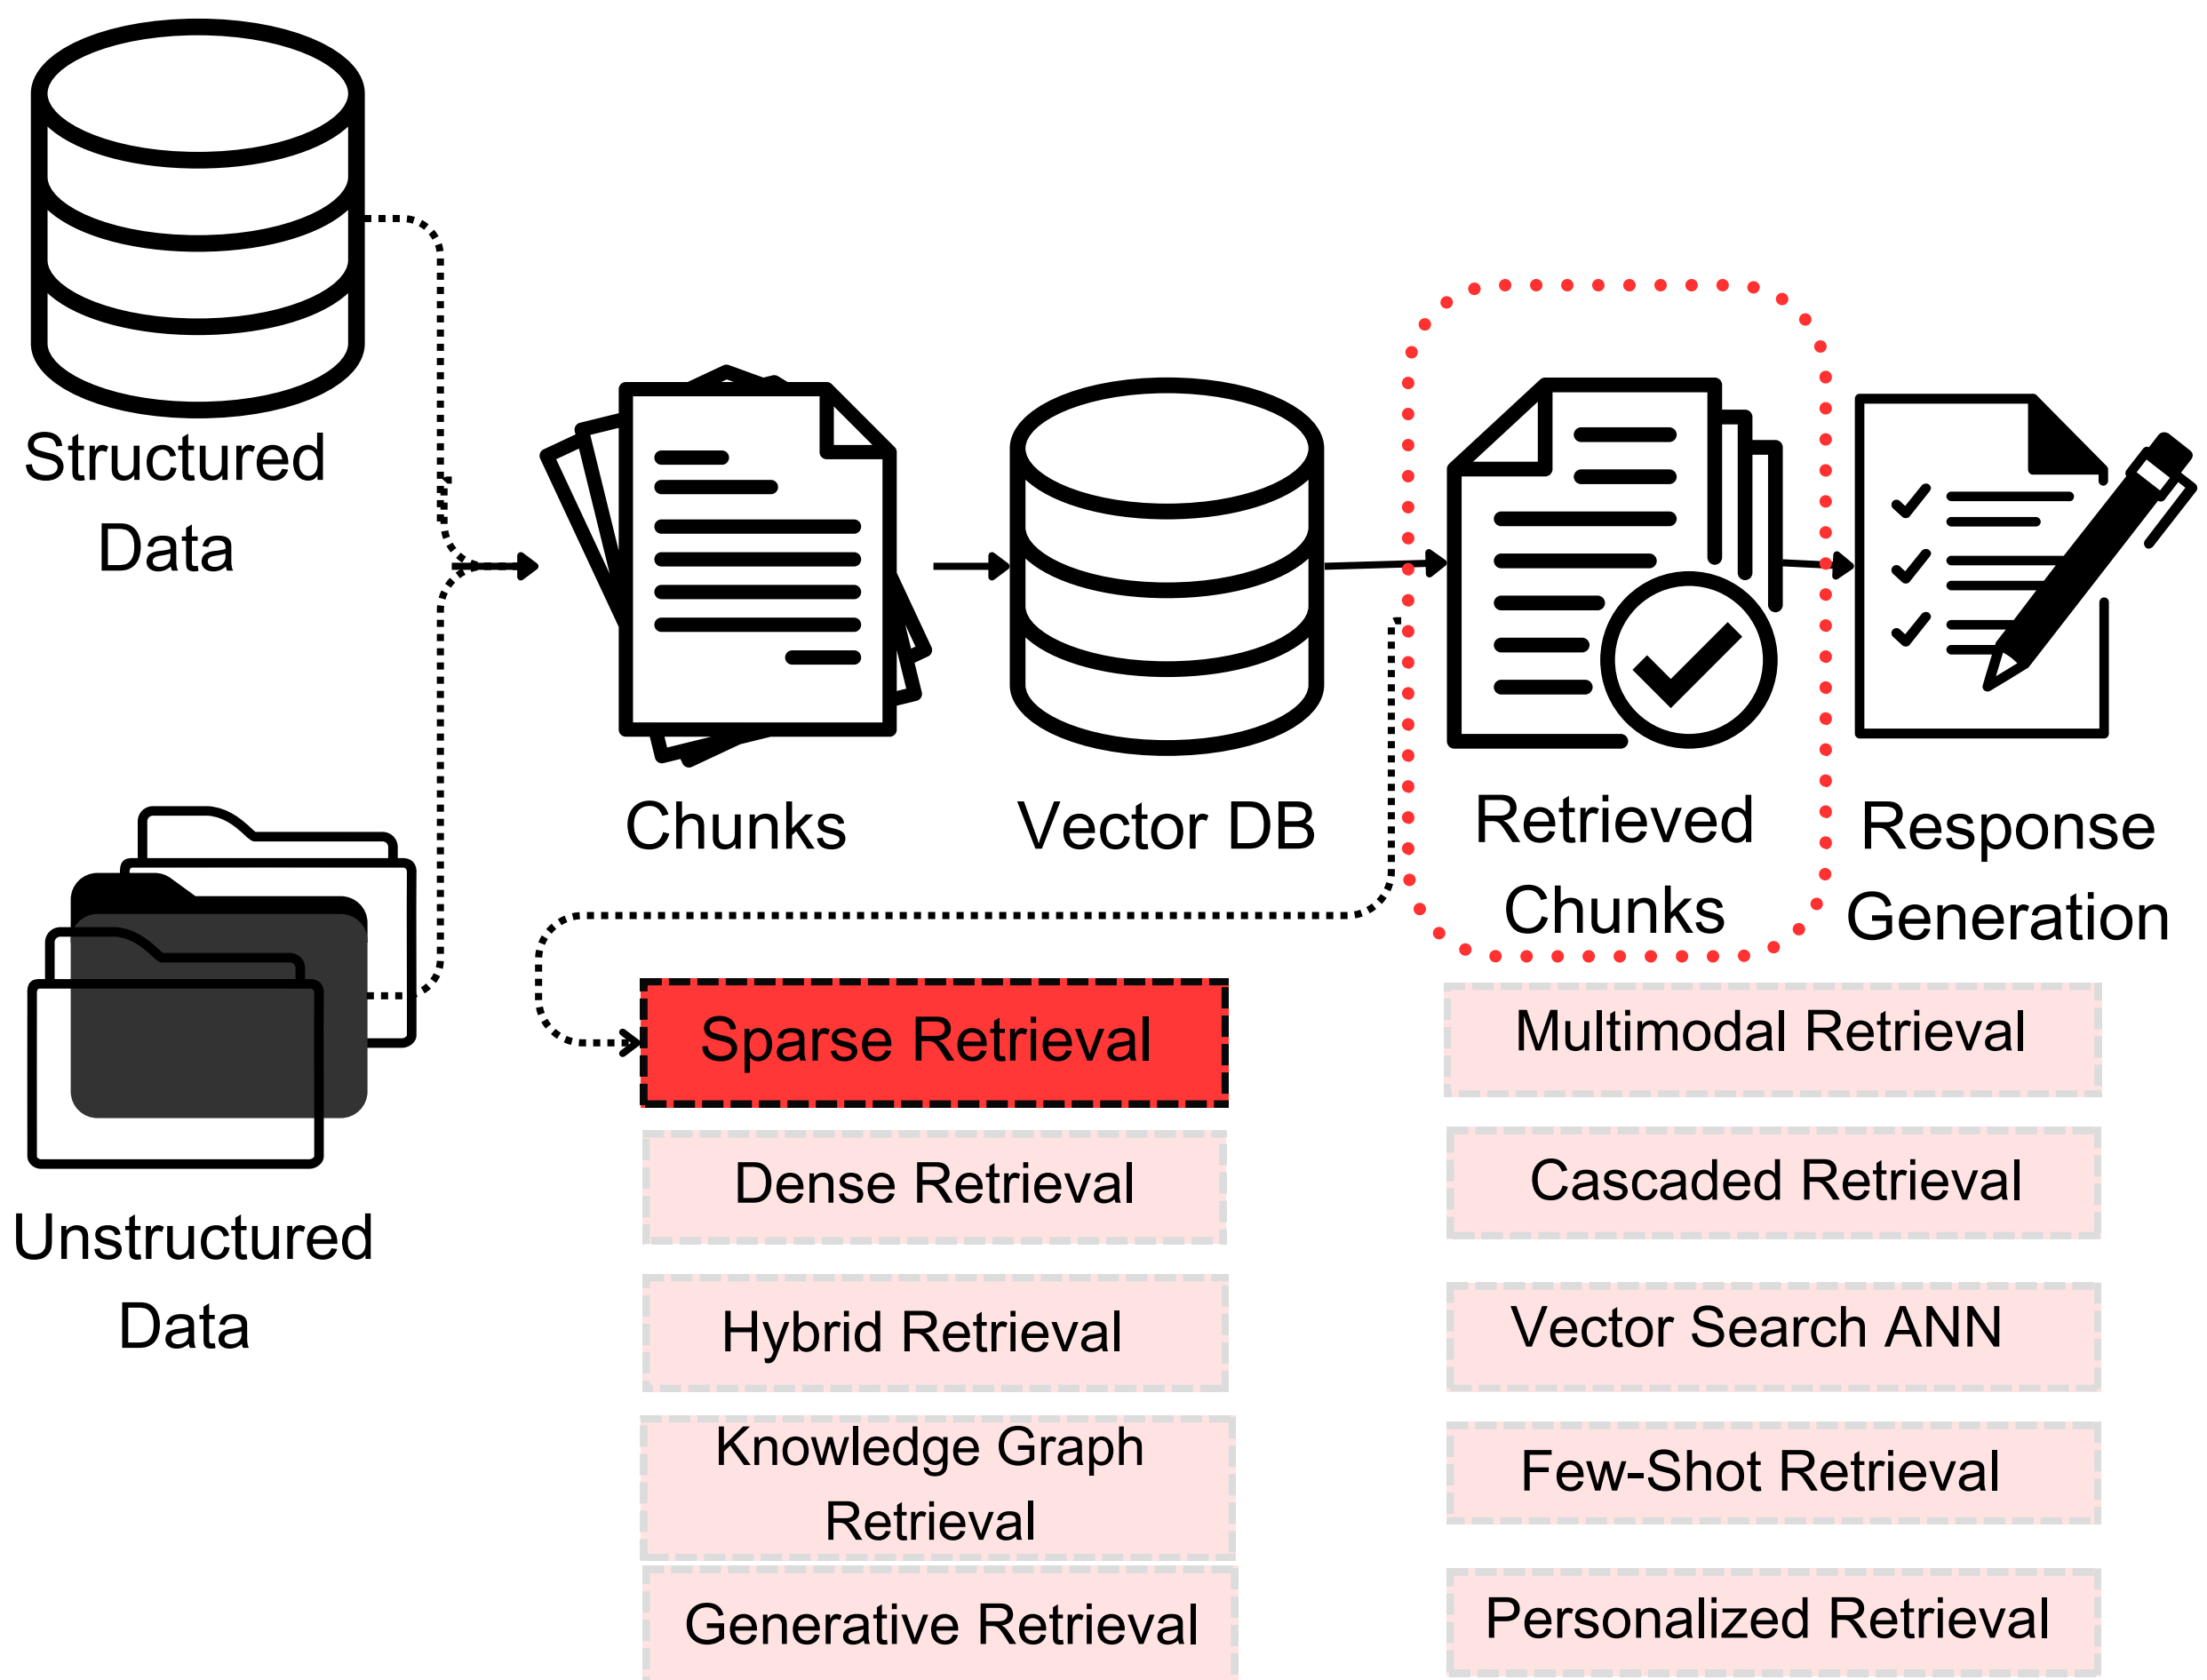



# Different Types of Retrieval in RAG System



# Information Retrieval



Information Retrieval (IR) is the process of finding relevant information from large collections of unstructured data like text, documents, or multimedia.

## How it works?

- 1 A user query is submitted (e.g., a question or keywords).
- 2 The system searches through available data (e.g., documents or embeddings).
- 3 Results are ranked and displayed based on relevance.

## Why is it important?

- 1 Powers RAG-based systems .
- 2 Enables context-aware conversational AI.
- 3 Supports knowledge-intensive tasks.

## Key Components


- 🔍 Retrieval Models: How the system retrieves data (e.g., Sparse, Dense).
- 📊 Ranking Algorithms: Ensuring the best results appear first.
- 🔄 Relevance Feedback: Continuous improvement based on user interaction.

# Sparse

# Retrieval



Relies on **exact keyword matching** to retrieve relevant documents.



## How it works?

- 1 Counts term frequency in the document.
- 2 Scores based on the importance of terms (e.g., rare terms have higher weight).
- 3 Returns documents that match the query terms exactly.

# Sparse

# Retrieval



## Pros

- ✓ Simple and interpretable.
- ✓ Low computational cost.
- ✓ Great for structured or small datasets.



## Cons


- ✗ Struggles with synonyms or semantic meaning (e.g., "car" ≠ "automobile").
- ✗ Limited to exact term overlap.

Examples


- Legal document searches.
- Library catalogs.
- FAQ-based retrieval systems.

Process


Query



Keyword  
Matching



Scoring



Ranked  
Results

# Sparse Retrieval - Best Practices

Optimize Preprocessing

Remove stop words and use stemming or lemmatization.

Choose the Right Algorithm

Use BM25 for weighted keyword matching over plain TF-IDF.

Use Hybrid Retrieval

Combine sparse (keyword matching) and dense (semantic search) for better performance.

## Technical Insights

### How TF-IDF scores work

Term Frequency  
(TF)

Counts the occurrences of a word in a document.

Inverse Document  
Frequency (IDF):

Penalizes common terms (e.g., "the", "is").

$$\text{TF-IDF}(t, d) = \text{TF}(t, d) \times \text{IDF}(t)$$

### How BM25 improves upon TF-IDF

- Includes term saturation to avoid over-penalizing long documents.
- Allows tuning with parameters like  $k_1$  (term frequency scaling) &  $b$  (length normalization).

# Sparse Retrieval in Python- TF-IDF

## Step1

### Step 1: Tokenization

First, we need to tokenize the documents.

```
documents = [
    "The cat sat on the mat",
    "The dog sat on the log",
    "The cat chased the dog"
]

# Tokenize the documents
tokenized_documents = [doc.lower().split() for doc in documents]
print(tokenized_documents)
```

```
[['the', 'cat', 'sat', 'on', 'the', 'mat'], ['the', 'dog', 'sat', 'on', 'the', 'log'], ['the', 'cat', 'chased', 'the', 'dog']]
```

## Step2

### Step 2: Calculate Term Frequency (TF)

Next, we calculate the term frequency for each term in each document.

```
from collections import Counter

# Calculate term frequency for each document
tf = [Counter(doc) for doc in tokenized_documents]
print(tf)
```

Python

```
[Counter({'the': 2, 'cat': 1, 'sat': 1, 'on': 1, 'mat': 1}), Counter({'the': 2, 'dog': 1, 'sat': 1, 'on': 1, 'log': 1}), Counter({'the': 1, 'cat': 1, 'chased': 1, 'the': 1, 'dog': 1})]
```

## Step3

### Step 3: Calculate Inverse Document Frequency (IDF)

Now, we calculate the inverse document frequency for each term across all documents.

```
import math

# Calculate document frequency for each term
df = Counter()
for doc in tokenized_documents:
    df.update(set(doc))

# Calculate IDF for each term
idf = {term: math.log(len(documents) / df[term]) for term in df}
print(idf)
```

```
{'on': 0.4054651081081644, 'cat': 0.4054651081081644, 'the': 0.0, 'sat': 0.4054651081081644, 'mat': 1.0986122886681098, 'dog': 0.4054651081081644, 'log': 0.4054651081081644, 'chased': 0.4054651081081644}
```



# Sparse Retrieval in Python - Tf-idf

## Step4

### Step 4: Calculate TF-IDF

Multiply TF by IDF for each term in each document.

```
# Calculate TF-IDF for each document
tf_idf = []
for doc_tf in tf:
    doc_tf_idf = {term: freq * idf[term] for term, freq in doc_tf.items()}
    tf_idf.append(doc_tf_idf)
print(tf_idf)
```

```
[{'the': 0.0, 'cat': 0.4054651081081644, 'sat': 0.4054651081081644, 'on': 0.4054651081081644, 'mat': 1.0986122886681098},
```

## Step5

### Step 5: Query Matching

```
# Tokenize the query
query = "cat sat"
tokenized_query = query.lower().split()

# Calculate term frequency for the query
query_tf = Counter(tokenized_query)

# Calculate TF-IDF for the query
query_tf_idf = {term: freq * idf.get(term, 0) for term, freq in query_tf.items()}
print(query_tf_idf)

# Calculate cosine similarity between query and each document
def cosine_similarity(doc_tf_idf, query_tf_idf):
    dot_product = sum(doc_tf_idf.get(term, 0) * query_tf_idf.get(term, 0) for term in query_tf_idf)
    doc_magnitude = math.sqrt(sum(value ** 2 for value in doc_tf_idf.values()))
    query_magnitude = math.sqrt(sum(value ** 2 for value in query_tf_idf.values()))
    if doc_magnitude == 0 or query_magnitude == 0:
        return 0.0
    return dot_product / (doc_magnitude * query_magnitude)

# Calculate similarity for each document
similarities = [cosine_similarity(doc_tf_idf, query_tf_idf) for doc_tf_idf in tf_idf]
print(similarities)

# Find the most relevant document
most_relevant_doc_index = similarities.index(max(similarities))
print(f"The most relevant document is: {documents[most_relevant_doc_index]}")
```

## Output

```
{'cat': 0.4054651081081644, 'sat': 0.4054651081081644}
[0.43976863279651823, 0.21988431639825912, 0.23135443112611218]
The most relevant document is: The cat sat on the mat
```

# Sparse Retrieval in Python - Tf-idf vs BM25

Tf-idf

BM25

Output

```
# Combine the above code into a single script
from sklearn.feature_extraction.text import TfidfVectorizer
import numpy as np
from rank_bm25 import BM25Okapi

corpus = [
    "The quick brown fox jumps over the lazy dog.",
    "Never jump over the lazy dog quickly.",
    "A quick brown dog outpaces a quick fox.",
    "Lazy dogs are not quick to jump over."
]

# TF-IDF Retrieval
tfidf_vectorizer = TfidfVectorizer()
tfidf_matrix = tfidf_vectorizer.fit_transform(corpus)
query = "quick fox"
query_vector = tfidf_vectorizer.transform([query])
cosine_similarities = np.dot(tfidf_matrix, query_vector.T).toarray().flatten()

print("TF-IDF Scores:")
for i, score in enumerate(cosine_similarities):
    print(f"Document {i+1}: {score}")

# BM25 Retrieval
tokenized_corpus = [doc.split(" ") for doc in corpus]
bm25 = BM25Okapi(tokenized_corpus)
tokenized_query = query.split(" ")
bm25_scores = bm25.get_scores(tokenized_query)

print("\nBM25 Scores:")
for i, score in enumerate(bm25_scores):
    print(f"Document {i+1}: {score}")
```

```
TF-IDF Scores:
Document 1: 0.402260298067896
Document 2: 0.0
Document 3: 0.684442866742917
Document 4: 0.16614037331333648
```

```
BM25 Scores:
Document 1: 0.9329649710463678
Document 2: 0.0
Document 3: 0.19735198611503196
Document 4: 0.13814639028052236
```



# CONGRATULATIONS

You have reached the end, now

If you want to help your network

REPOST THIS




Sarveshwaran R


<https://github.com/DataSphereX/Retrieval-Strategies>




**DataSphereX/Retrieval-Strategies**

Retrieval Strategies for RAG




 1


Contributor

 0


Issues

 0

Stars


 0

Forks



**DataSphereX/Retrieval-Strategies: Retrieval Strategies for RAG**

Retrieval Strategies for RAG . Contribute to DataSphereX/Retrieval-Strategies development by creating an account on GitHub.

 GitHub